

# Grocery Sales Forecast

Christine Vu<sup>1</sup>, Dave Friesen<sup>2</sup>

12/12/2022

## Get Data

Reference: [<https://www.kaggle.com/competitions/store-sales-time-series-forecasting/data?select=transactions.csv>  
(<https://www.kaggle.com/competitions/store-sales-time-series-forecasting/data?select=transactions.csv>)]

```
#dir <- "/Users/christinevu/Downloads/"
dir <- "../data/"

# Load dataset(s)
setwd(dir)
sales_df <- read.csv("train.csv", header = TRUE)
sales_test_df <- read.csv("test.csv", header = TRUE)
stores_df <- read.csv("stores.csv", header = TRUE)
oil_df <- read.csv("oil.csv", header = TRUE)
events_df <- read.csv("holidays_events.csv", header = TRUE)

# Data validation and understanding, including structure, content, and statistical characteristics covered below
```

## Explore & Visualize Series

### Univariate Analysis and Preliminary Pre-Processing

```
# e.g., statistical characteristics (including distribution, skewness, outliers)
# +[optionally] review sample observations
univariate(sales_df); head(sales_df, 3); #str(sales_df)
```

Summary Univariate Analysis for (sales_df) (3,000,888 observations)										
	Type	NA%	Blank%	Unique Freq	Min	Max	Mean	Median	Outlier<	>Outlier
id	integer			3000888		3000887	1500444	No	No	N 1.0
date	character			1684	1					N 1.0
store_nbr	integer			54	1	54	28	No	Yes	N 1.0
family	character			33						N 0.8
sales	numeric	31%		379610		124717.0	357.8	11.0	No	Yes 7.4 N
onpromotion	integer	79%		362		741		No	No	11.2 N

	id	date	store_nbr	family	sales	onpromotion
1	0	2013-01-01	1	AUTOMOTIVE	0	0
2	1	2013-01-01	1	BABY CARE	0	0
3	2	2013-01-01	1	BEAUTY	0	0

```
univariate(sales_test_df); head(sales_test_df, 3); #str(sales_test_df)
```

Summary Univariate Analysis for (sales_test_df) (28,512 observations)										
	Type	NA%	Blank%	Unique Freq	Min	Max	Mean	Median	Outlier<	>Outlier
id	integer			28512	3000888	3029399	3015144	No	No	N 1.0
date	character			16	1					N 1.0
store_nbr	integer			54	1	54	28	No	Yes	N 1.0
family	character			33						N 0.8
onpromotion	integer	55%		212		646		No	Yes	8.5 N

	id	date	store_nbr	family	onpromotion
1	3000888	2017-08-16	1	AUTOMOTIVE	0
2	3000889	2017-08-16	1	BABY CARE	0
3	3000890	2017-08-16	1	BEAUTY	2

```
univariate(stores_df); head(stores_df, 3); #str(stores_df)
```

```
Summary Univariate Analysis for (stores_df) (54 observations)
      Type      NA% Blank% Unique Freq     Min     Max     Mean     Median Outlier< >Outlier Skewness nZV ACF1
store_nbr integer          54             1       54        28      No      No           N   0.9
city      character         22            16           5           N   0.3
state      character         16            5           N   0.4
type      character          5            17           1           N   0.6
cluster    integer           17            8           Yes          N   0.2
```

	store_nbr	city	state	type	cluster
1	1	Quito	Pichincha	D	13
2	2	Quito	Pichincha	D	13
3	3	Quito	Pichincha	D	8

```
univariate(oil_df); head(oil_df, 3); #str(oil_df)
```

```
Summary Univariate Analysis for (oil_df) (1,218 observations)
      Type      NA% Blank% Unique Freq     Min     Max     Mean     Median Outlier< >Outlier Skewness nZV ACF1
date      character         1218           1           3           N   1.0
dcoilwtico numeric          998            26.2        110.6      67.7      53.2      No      Yes           0.3   N   0.9
```

	date	dcoilwtico
1	2013-01-01	NA
2	2013-01-02	93
3	2013-01-03	93

```
univariate(events_df); head(events_df, 3); #str(events_df)
```

```
Summary Univariate Analysis for (events_df) (350 observations)
      Type      NA% Blank% Unique Freq     Min     Max     Mean     Median Outlier< >Outlier Skewness nZV ACF1
date      character         312            7           3           N   1.0
type      character          6            24           1           N   0.3
locale    character          3            103          2           N   0.3
locale_name character         24           103          2           N   0.2
description character         103           2           Y           N   0.2
transferred character          2            1           Y           N   0.2
```

	date	type	locale	locale_name	description	transferred
1	2012-03-02	Holiday	Local	Manta	Fundacion de Manta	False
2	2012-04-01	Holiday	Regional	Cotopaxi	Provincializacion de Cotopaxi	False
3	2012-04-12	Holiday	Local	Cuenca	Fundacion de Cuenca	False

```
# The following code was key for preliminary EDA, with individual results ultimately summarized
# through the above univariate() calls
```

```
# Data summary statistics
#summary(sales_df)
#summary(sales_test_df)
#summary(stores_df)
#summary(oil_df)
#summary(events_df)

# Data types
#lapply(sales_df, class)
#lapply(sales_test_df, class)
#lapply(stores_df, class)
#lapply(oil_df, class)
#lapply(events_df, class)

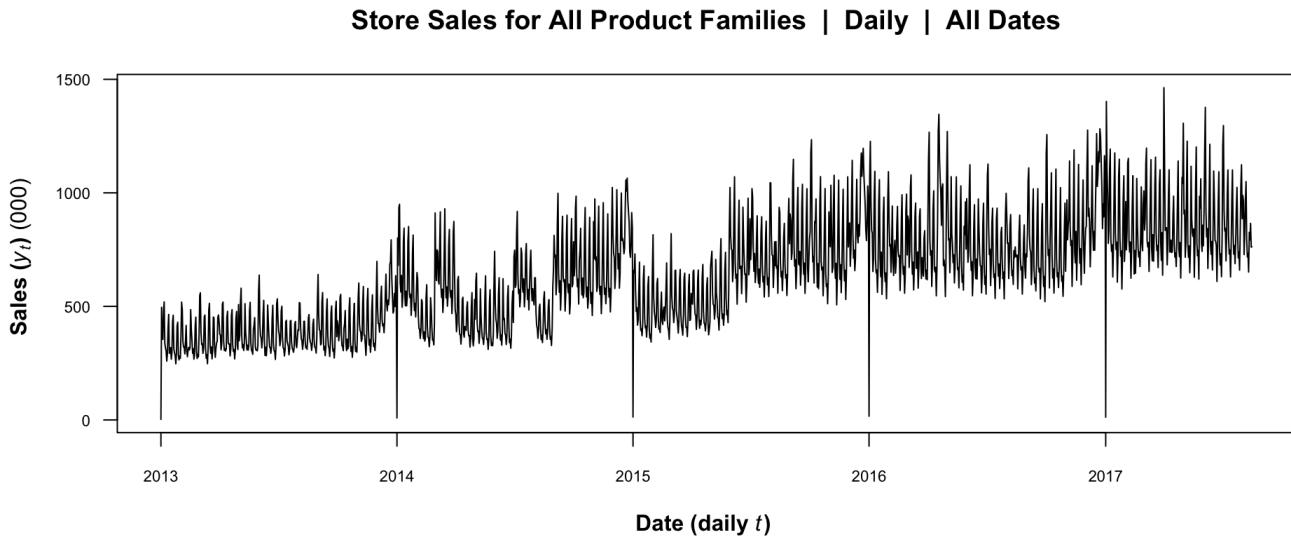
# Row and column counts
#nrow(sales_df)
#row(sales_test_df)
#ncol(sales_df)
#ncol(sales_test_df)

# Missing values
#sum(is.na(sales_df))
#sum(is.na(sales_test_df))
#sum(is.na(stores_df))
#sum(is.na(oil_df))
#sum(is.na(events_df))
```

```
# Convert string mm/dd/yyyy to Date values and confirm sort
sales_df <- (sales_df %>%
  mutate(date = as.Date(date, format = "%Y-%m-%d"),) %>%
  arrange(date))
```

## Series Visualization - All Product Families

```
# Aggregate base dataframe by date (i.e., sum all product lines) and initially plot series at
# provided time granularity
sales_agg_df <- as.data.frame(sales_df %>%
  group_by(date) %>%
  summarize(sales = sum(sales / 1000.0)))
plot(x = sales_agg_df$date, y = sales_agg_df$sales, type = "l",
  main = "Store Sales for All Product Families | Daily | All Dates",
  xlab = TeX(r"(\textbf{Date (daily $t$)})"), ylab = TeX(r"(\textbf{Sales (\textit{$y_t$})} (000))"),
  las = 1, cex.axis = 0.7)
```



```
# Aggregate base dataframe from daily to weekly for all product families
sales_wk_df <- as.data.frame(sales_df %>%
  mutate(year = year(date), week = week(date)) %>%
  group_by(year, week) %>%
  summarize(sales = sum(sales / 1000.0)))

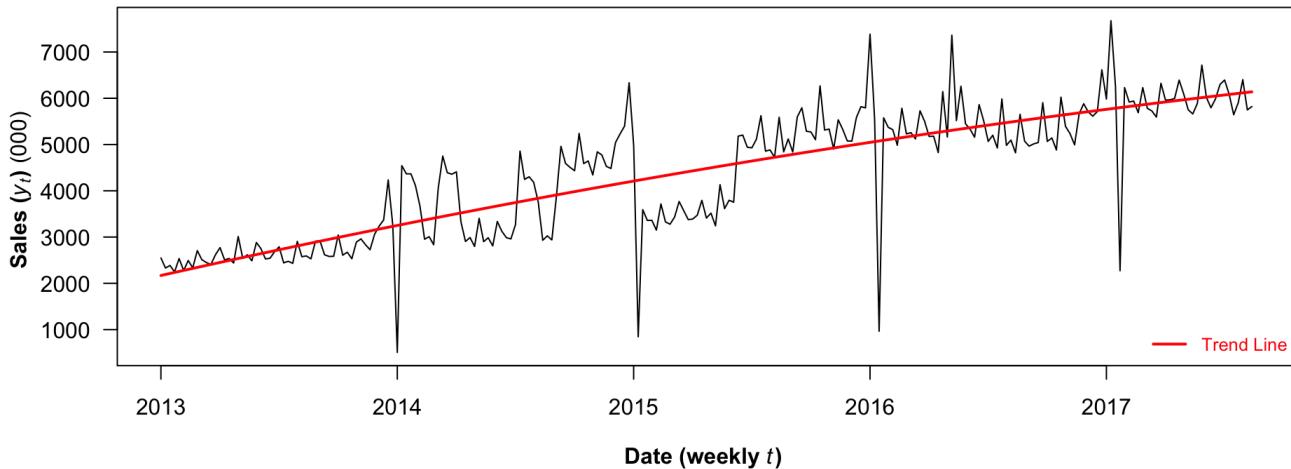
# Create overall time series
sales_begin_year <- head(sales_wk_df$year, 1)
sales_begin_week <- head(sales_wk_df$week, 1)
sales_end_year <- tail(sales_wk_df$year, 1)
sales_end_week <- tail(sales_wk_df$week, 1)
sales_ts <- ts(sales_wk_df$sales,
  start = c(sales_begin_year, sales_begin_week),
  end = c(sales_end_year, sales_end_week), freq = 52)

# Plot overall time series with trend line
plot(sales_ts, type = "l",
  main = "Store Sales for All Product Families | Weekly | All Dates",
  xlab = TeX(r"(\textbf{Date (weekly \textit{$t$})})"), ylab = TeX(r"(\textbf{Sales (\textit{$y_t$})} (000))"),
  las = 1, cex.axis = 0.7)

sales_lm <- tslm(sales_ts ~ trend + I(trend^2))
lines(sales_lm$fitted, lwd = 2, lty = 1, col = "red")

legend("bottomright",
  legend = c("Trend Line"),
  col = c("red"),
  lwd = 2, lty = 1.2, cex = 0.8,
  box.lty = 0, bg = NULL, text.col = "red")
```

## Store Sales for All Product Families | Weekly | All Dates

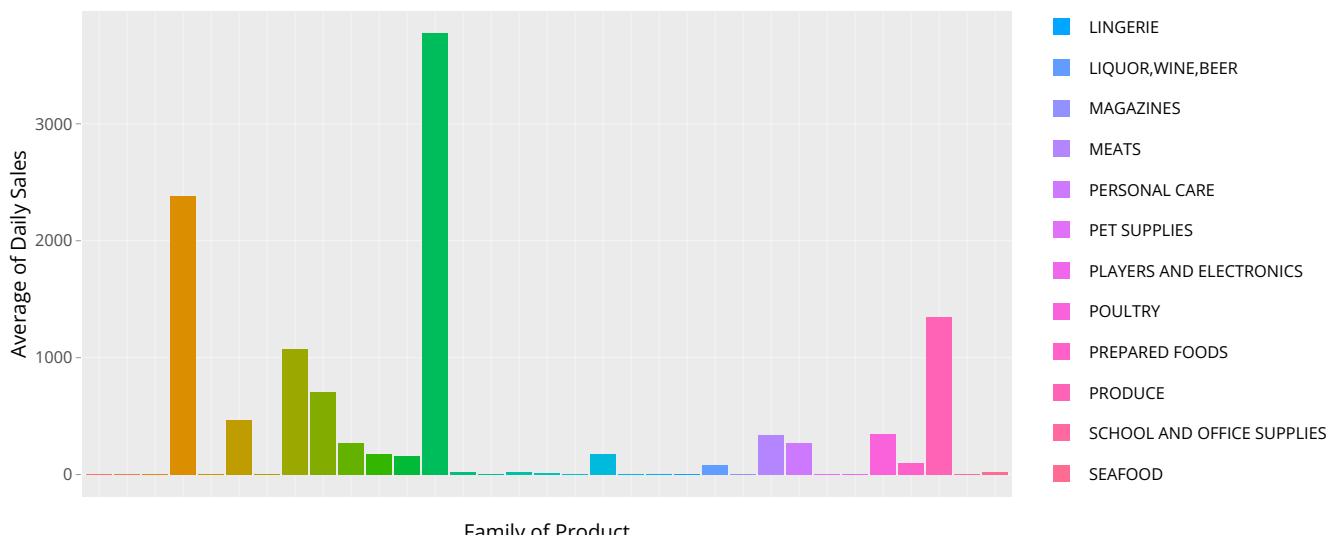


```
# Proportion by family
fam_prop <- sales_df %>%
  group_by(family) %>%
  summarise(Mean_Sales = round(mean(sales, na.rm = TRUE), 2),
            Median_Sales = round(median(sales, na.rm = TRUE), 2),
            Maximum_Sales = max(sales),
            IQR_Sales = IQR(sales),
            Total_Sales = sum(sales))
head(fam_prop, 3)
```

family	Mean_Sales	Median_Sales	Maximum_Sales	IQR_Sales	Total_Sales
AUTOMOTIVE	6.1	5	255	6	554822
BABY CARE	0.11	0	116	0	10051
BEAUTY	3.72	2	136	5	337893

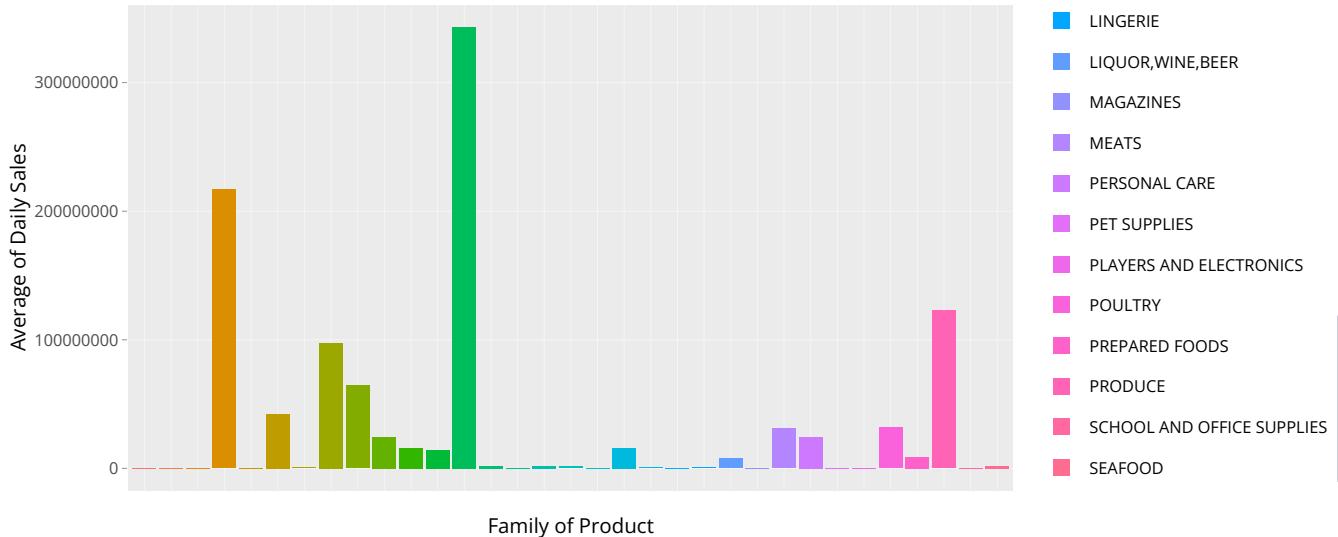
```
# Box plot: Mean of store sales by family
fam_plot <- ggplot(fam_prop) +
  geom_col(aes(x = family, y = Mean_Sales, fill = family)) +
  theme(axis.ticks.x = element_blank(),
        axis.text.x = element_blank()) +
  labs(title = "Average Store Sales by Family",
       x = "Family of Product",
       y = "Average of Daily Sales",
       legend = "Family of Product Sold",
       bty = "l")
ggplotly(fam_plot)
```

Average Store Sales by Family



```
# Box plot: Total store sales by family
fam_plot2 <- ggplot(fam_prop) +
  geom_col(aes(x = family, y = Total_Sales, fill = family)) +
  theme(axis.ticks.x = element_blank(),
        axis.text.x = element_blank()) +
  labs(title = "Total Store Sales by Family",
       x = "Family of Product",
       y = "Average of Daily Sales",
       legend = "Family of Product Sold",
       bty = "l")
ggplotly(fam_plot2)
```

Total Store Sales by Family



```
# Extract day, month, year, etc. from sales_df
sales_df$day <- sales_df$date %>% day()
sales_df$month <- sales_df$date %>% month()
sales_df$month_lab <- sales_df$date %>% month(label = TRUE)
sales_df$year <- sales_df$date %>% year()
sales_df$week <- sales_df$date %>% week()
sales_df$week_day <- sales_df$date %>%
  wday(week_start = getOption("lubridate.week.start", 1), label = TRUE)
```

```
opar = par()
par(mfrow = c(1, 2))

# Plot overall time series w/log scale
plot(sales_ts, type = "l",
      main = "Store Sales | Weekly | Log Scale",
      xlab = TeX(r"(\textbf{Date (weekly \textit{$t$})})"),
      ylab = TeX(r"(\textbf{Sales (\textit{$y_t$})} (000))"),
      las = 1, cex.axis = 0.7,
      log = "y")

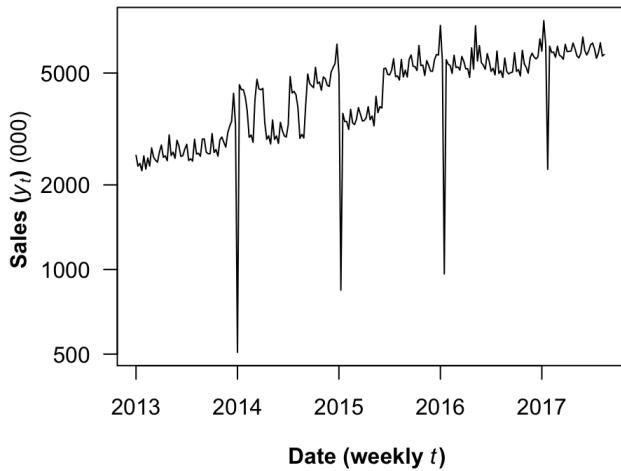
# Plot zoomed time series with trend line
sales_zoom_ts <- window(sales_ts, start = c(sales_end_year, 1), end = c(sales_end_year, 13))
plot(sales_zoom_ts, type = "l",
      main = "Store Sales | Weekly | One Quarter",
      xlab = TeX(r"(\textbf{Date (weekly \textit{$t$})})"),
      ylab = TeX(r"(\textbf{Sales (\textit{$y_t$})} (000))"),
      xaxt = "n",
      las = 1, cex = 0.7)

sales_zoom_lm <- tslm(sales_zoom_ts ~ trend + I(trend^2))
lines(sales_zoom_lm$fitted, lwd = 2, lty = 1, col = "red")

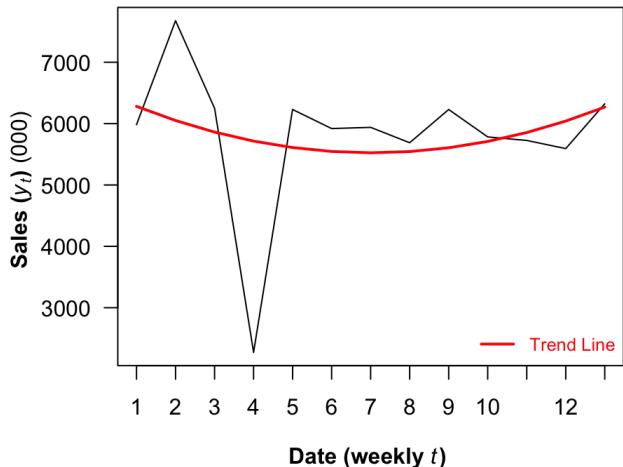
axis(1, at = as.numeric(time(sales_zoom_ts)), labels = seq(sales_zoom_ts))

legend("bottomright",
       legend = "Trend Line",
       col = "red",
       lwd = 2, lty = 1.2, cex = 0.8,
       box.lty = 0, bg = NULL, text.col = "red")
```

**Store Sales | Weekly | Log Scale**



**Store Sales | Weekly | One Quarter**



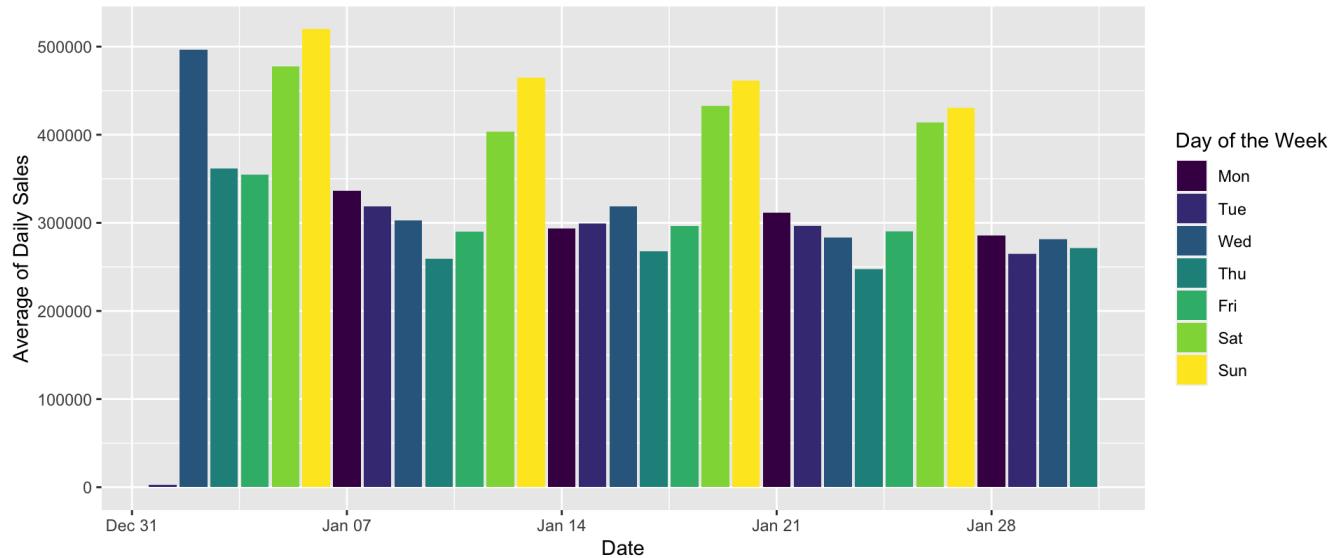
```
par(opar)
```

```
# Proportion by sales
sales_prop <- sales_df %>%
  group_by(store_nbr) %>%
  summarise(Mean_Sales = mean(sales),
            Median_Sales = median(sales),
            Maximum_Sales = max(sales),
            IQR_Sales = IQR(sales),
            IQR_to_avg_sales = IQR(sales) / mean(sales),
            Total_Sales = sum(sales))
head(sales_prop, 3)
```

store_nbr	Mean_Sales	Median_Sales	Maximum_Sales	IQR_Sales	IQR_to_avg_sales	Total_Sales
1	255.	19	9065	165	0.648	14145013.
2	388.	23.2	124717	289	0.745	21557389.
3	908.	64	21858	582	0.641	50481910.

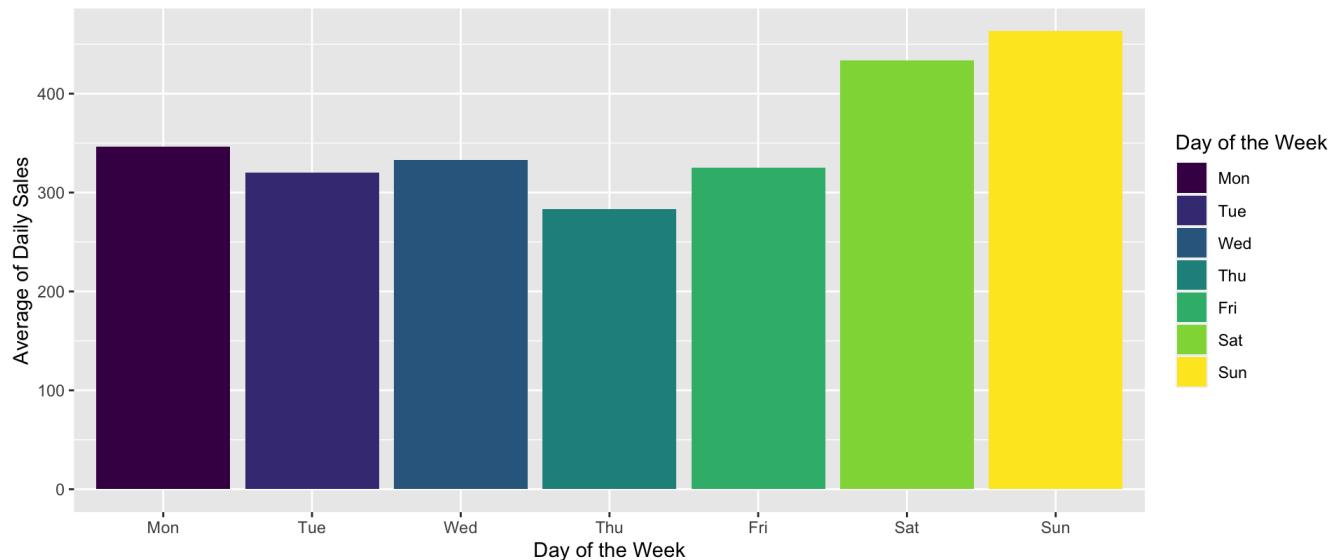
```
# Bar plot: Average of store sales by day of the week in January 2013
sales_plot <- sales_df %>%
  group_by(date) %>%
  summarise(avg_sales = mean(sales), wday = week_day, .groups = "keep") %>%
  filter(date <= '2013-01-31') %>%
  ggplot() +
  geom_col(aes(x = date, y = avg_sales, fill = wday)) +
  labs(title = "Average of Daily Store Sales in January 2013",
       x = "Date",
       y = "Average of Daily Sales",
       fill = "Day of the Week",
       bty = "l")
sales_plot
```

### Average of Daily Store Sales in January 2013

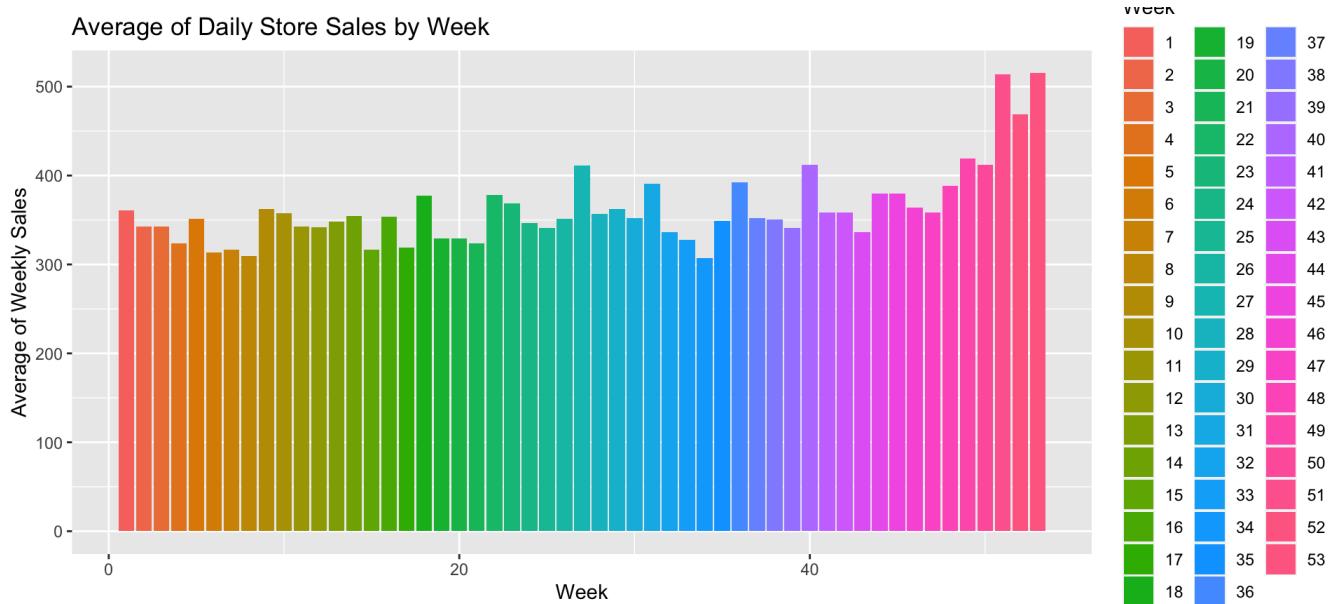


```
# Bar plot: Average of store sales by day of the week
sales_plot1 <- sales_df %>%
  group_by(week_day) %>%
  summarise(avg_sales = mean(sales)) %>%
  ggplot() +
  geom_col(aes(x = week_day,
                y = avg_sales,
                fill = week_day)) +
  labs(title = "Average of Daily Store Sales by Day of the Week",
       x = "Day of the Week",
       y = "Average of Daily Sales",
       fill = "Day of the Week",
       bty = "l")
sales_plot1
```

### Average of Daily Store Sales by Day of the Week

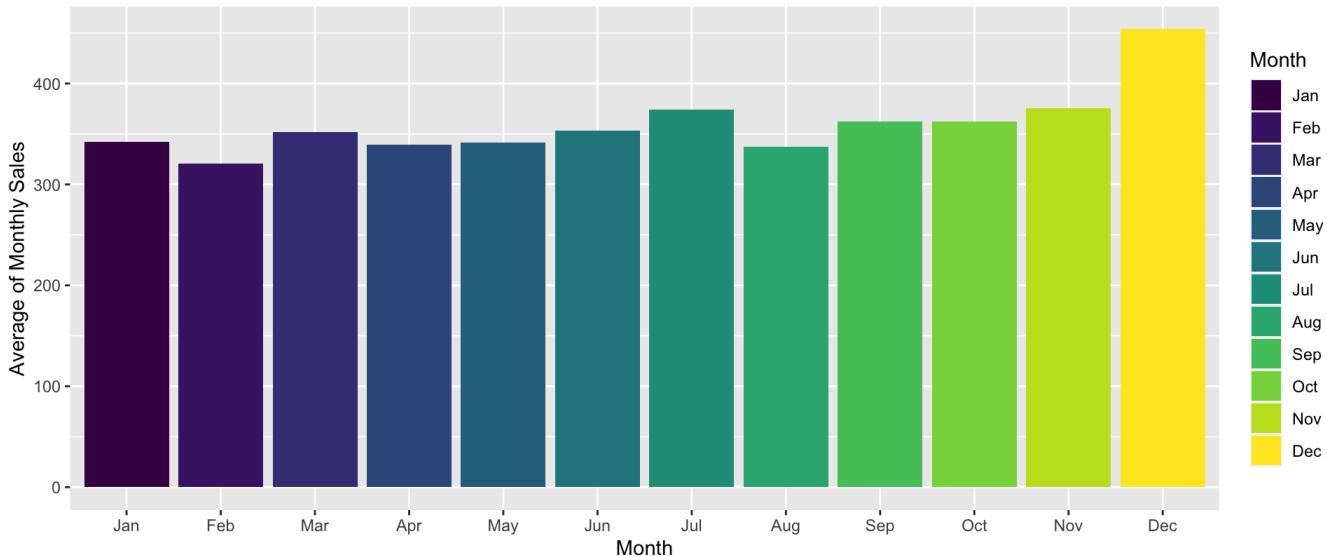


```
# Bar plot: Average of store sales by week
sales_plot2 <- sales_df %>%
  group_by(week) %>%
  summarise(avg_sales = mean(sales)) %>%
  ggplot() +
  geom_col(aes(x = week,
               y = avg_sales,
               fill = as.factor(week))) +
  labs(title = "Average of Daily Store Sales by Week",
       x = "Week",
       y = "Average of Weekly Sales",
       fill = "Week",
       bty = "l")
sales_plot2
```



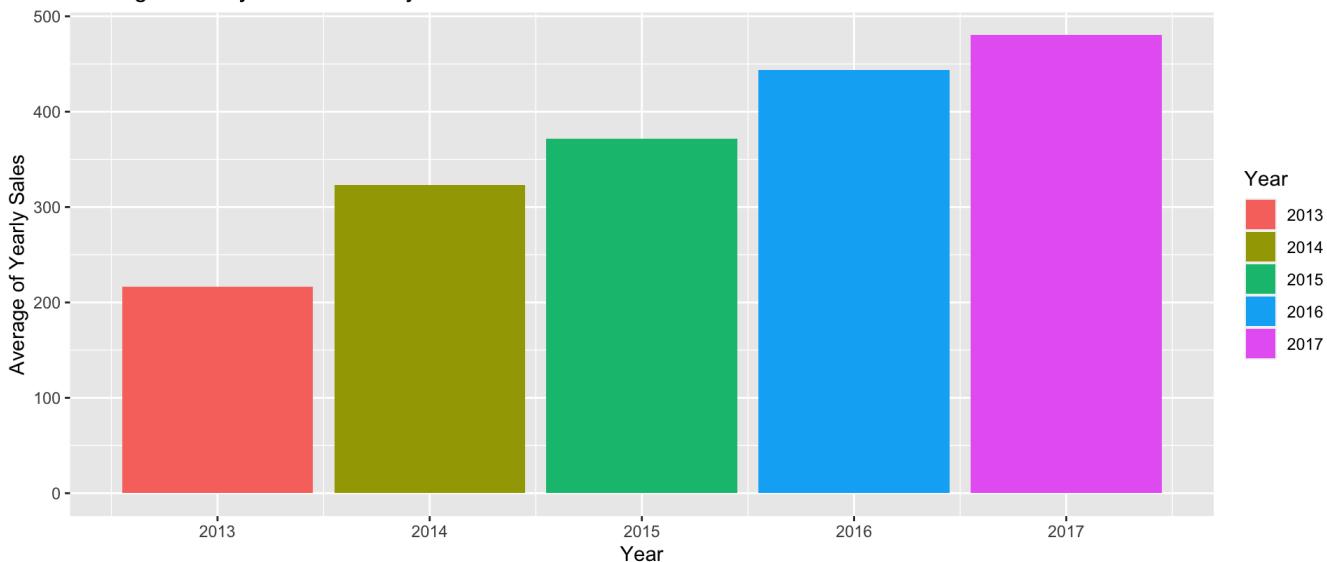
```
# Bar plot: Average of store sales by month of the year
sales_plot3 <- sales_df %>%
  group_by(month_lab) %>%
  summarise(avg_sales = mean(sales)) %>%
  ggplot() +
  geom_col(aes(x = month_lab,
               y = avg_sales,
               fill = month_lab)) +
  labs(title = "Average of Daily Store Sales by Month of the Year",
       x = "Month",
       y = "Average of Monthly Sales",
       fill = "Month",
       bty = "l")
sales_plot3
```

Average of Daily Store Sales by Month of the Year



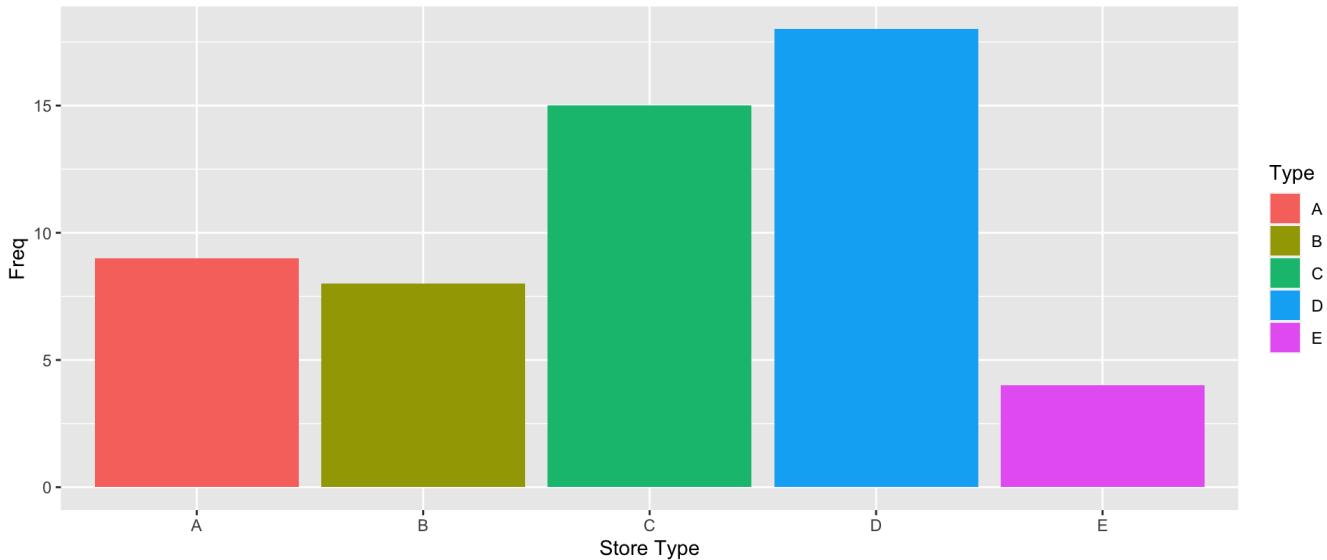
```
# Bar plot: Average of store sales by year
sales_plot4 <- sales_df %>%
  group_by(year) %>%
  summarise(avg_sales = mean(sales)) %>%
  ggplot() +
  geom_col(aes(x = year,
               y = avg_sales,
               fill = as.factor(year))) +
  labs(title = "Average of Daily Store Sales by Year",
       x = "Year",
       y = "Average of Yearly Sales",
       fill = "Year",
       bty = "l")
sales_plot4
```

Average of Daily Store Sales by Year



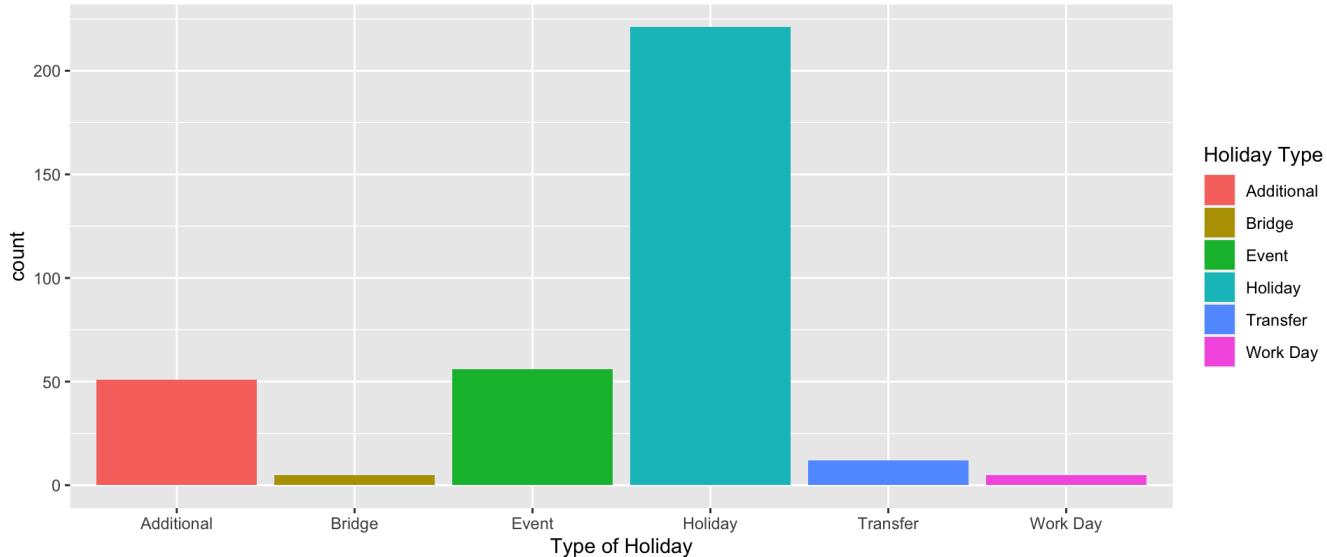
```
# Bar plot: Distribution of store types
store_plot <- stores_df$type %>%
  table() %>%
  as.data.frame() %>%
  ggplot() +
  geom_col(aes(y = Freq, x = ., fill = as.factor(.))) +
  labs(title = "Distribution of Store Types",
       x = "Store Type",
       fill = "Type",
       bty = "l")
store_plot
```

### Distribution of Store Types



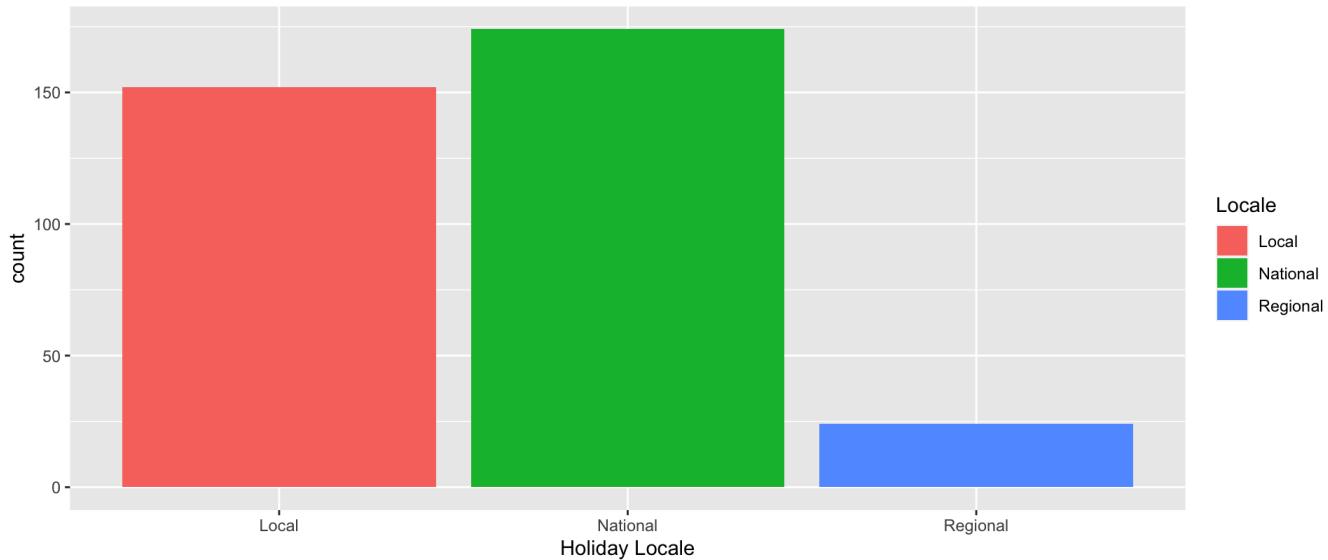
```
# Bar plot: Distribution of holiday type
events_plot1 <- events_df %>%
  ggplot() +
  geom_bar(aes(x = type, fill = type)) +
  labs(title = "Distrbution of Type of Holiday",
       x = "Type of Holiday",
       fill = "Holiday Type",
       bty = "l")
events_plot1
```

### Distrbution of Type of Holiday

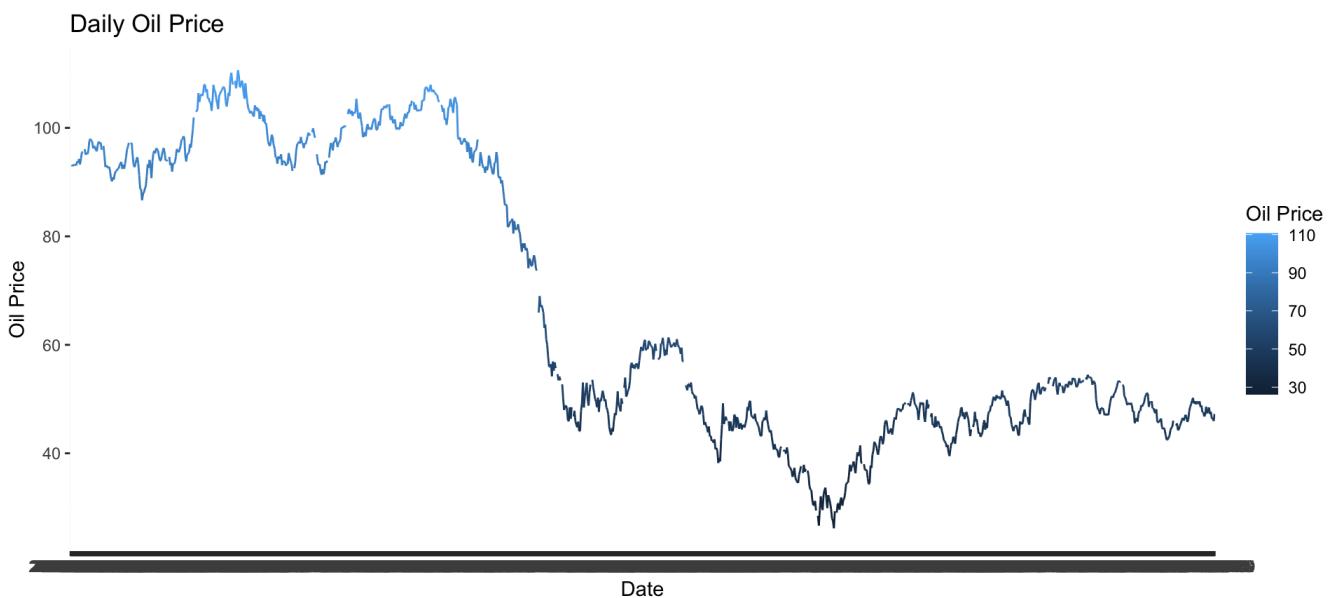


```
# Bar plot: Distribution of holiday locale
events_plot2 <- events_df %>%
  ggplot() +
  geom_bar(aes(x = locale, fill = locale)) +
  labs(title = "Distribution of Holiday Locale",
       x = "Holiday Locale",
       fill = "Locale",
       bty = "l")
events_plot2
```

### Distribution of Holiday Locale



```
# Line chart: Daily oil price
oil_plot <- oil_df %>%
  ggplot(aes(x = date, y = dcoilwtico, color = dcoilwtico, group = 1)) +
  geom_line(na.rm = TRUE) +
  labs(title = "Daily Oil Price",
       x = "Date",
       y = "Oil Price",
       color = "Oil Price",
       bty = "l")
oil_plot
```



### Series Visualization - Individual Product Families

```

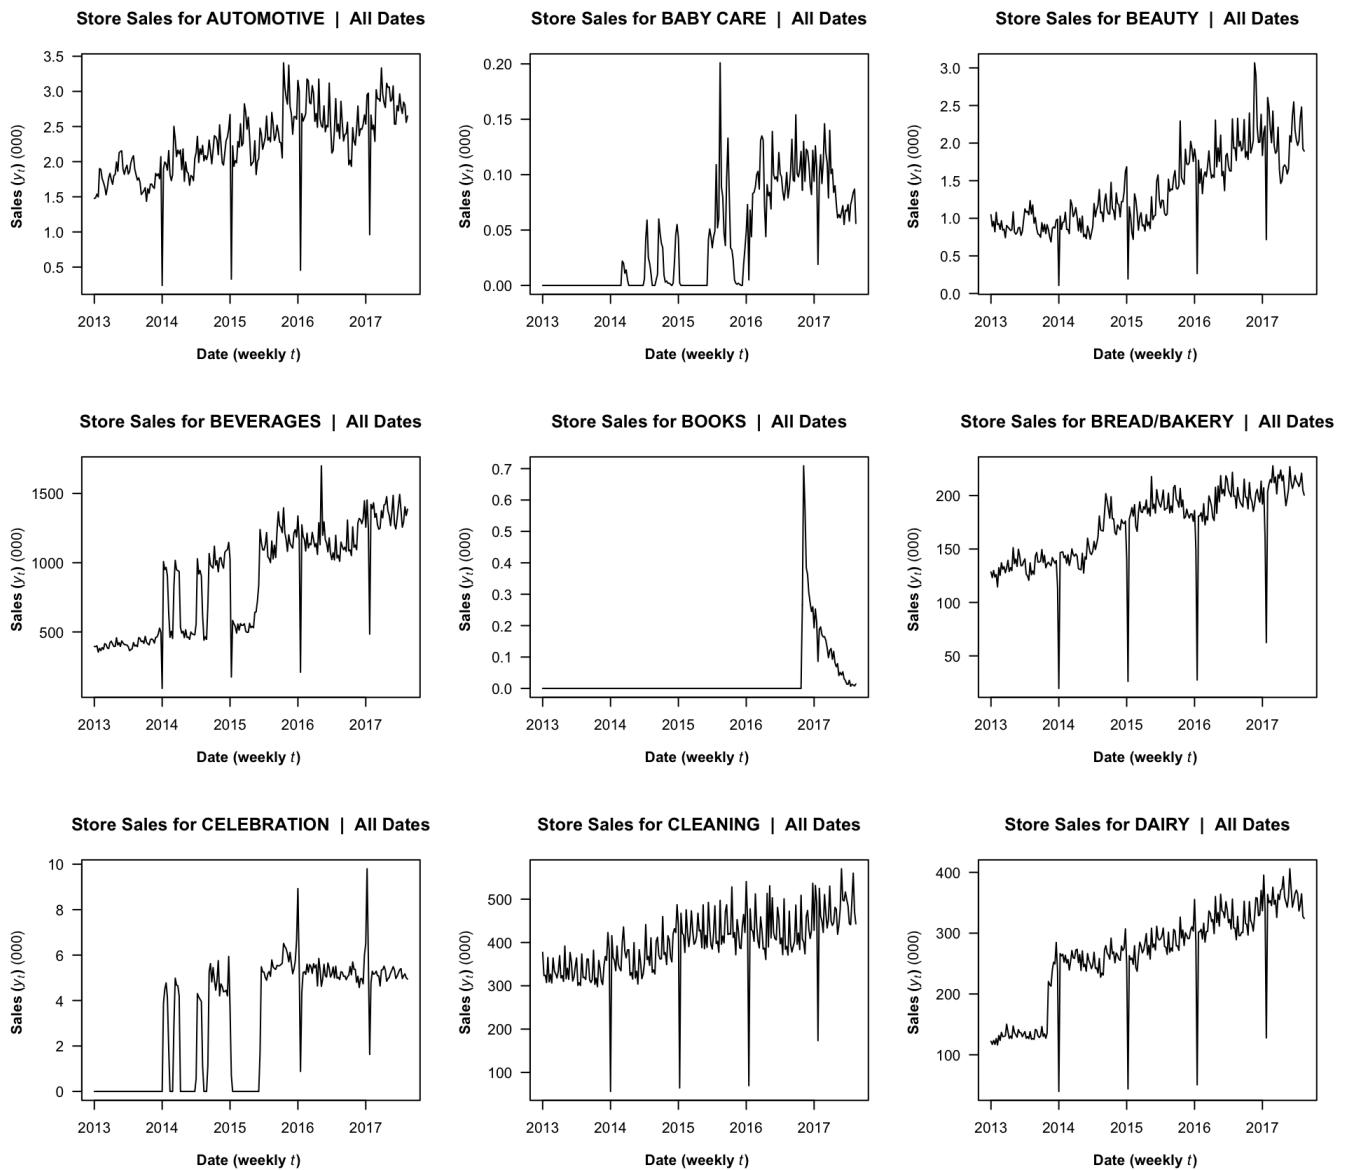
# Aggregate base dataframe from daily to weekly by product family
sales_wk_df <- as.data.frame(sales_df %>%
  mutate(year = year(date), week = week(date)) %>%
  group_by(family, year, week) %>%
  summarize(sales = sum(sales / 1000.0)))

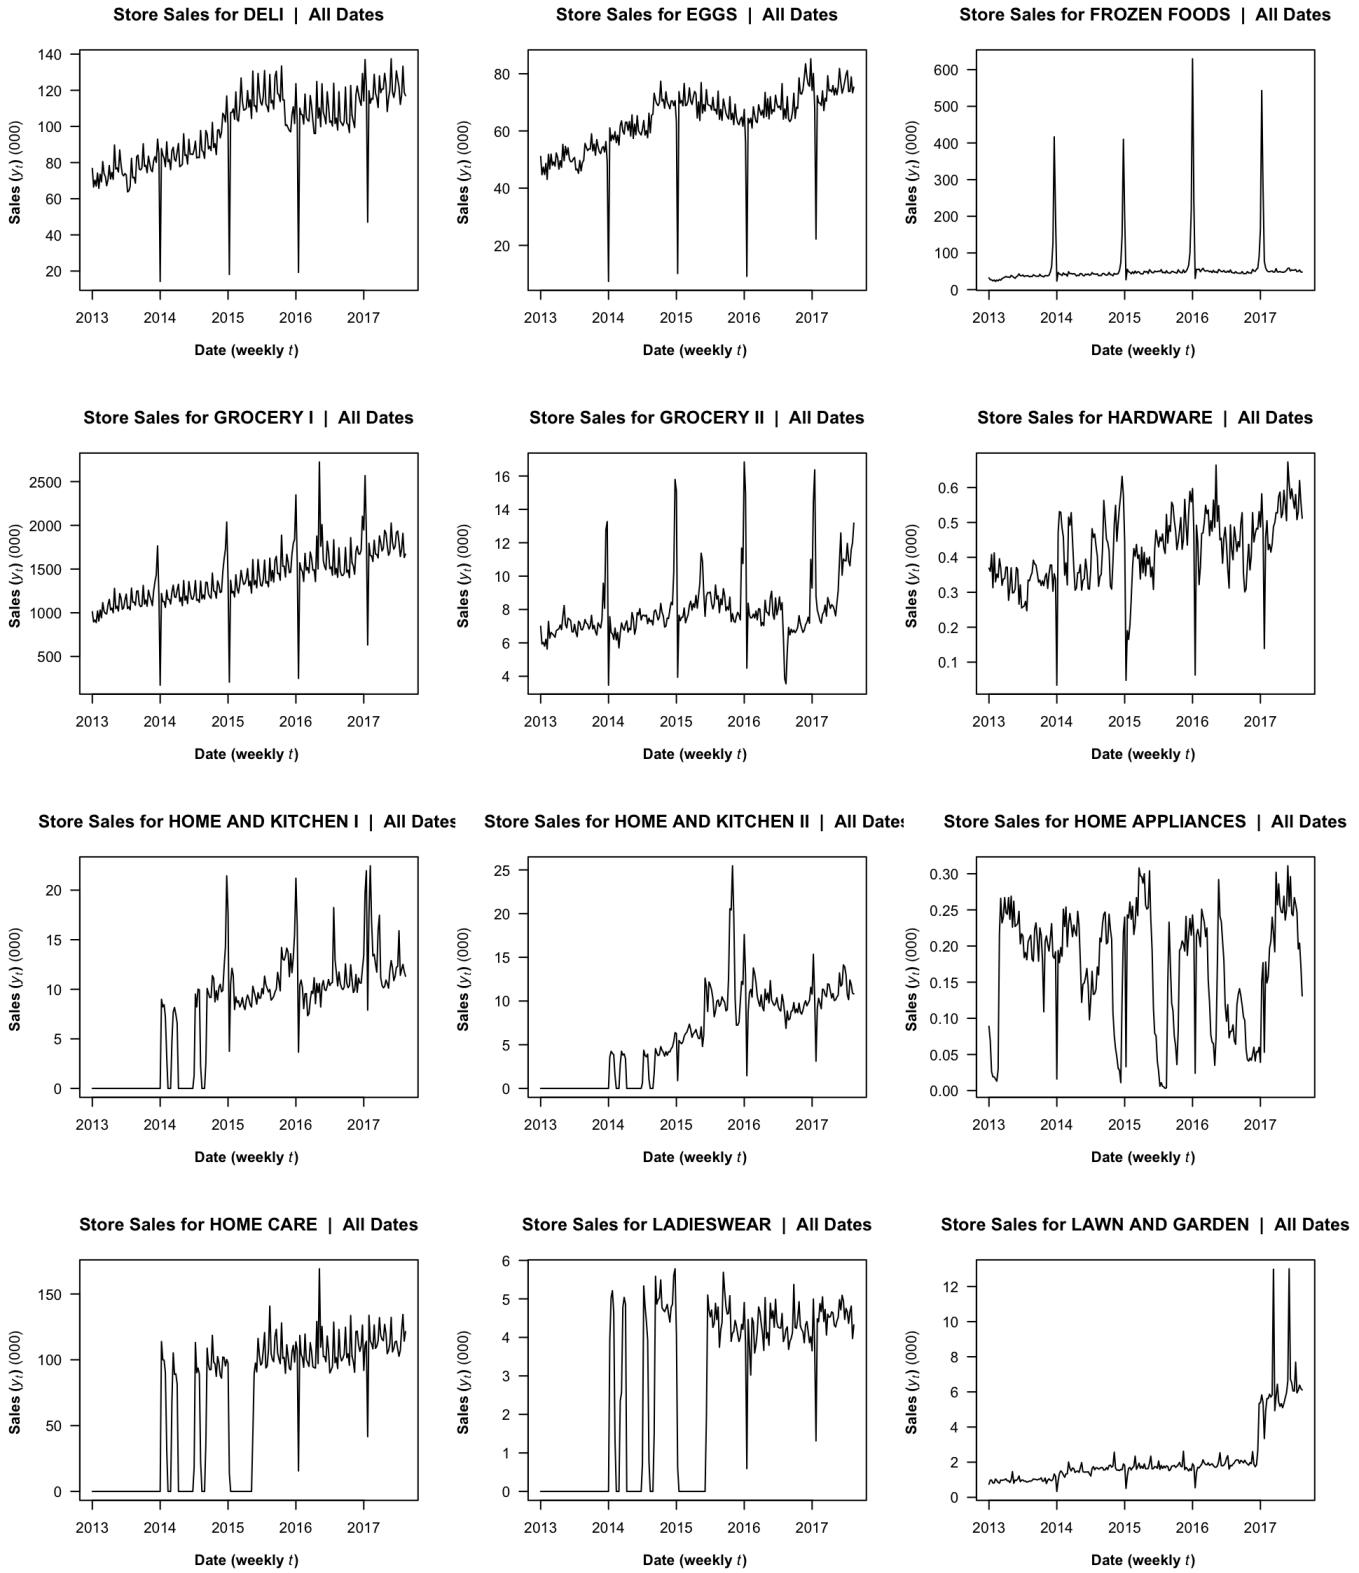
opar = par()
par(mfrow = c(1, 3))

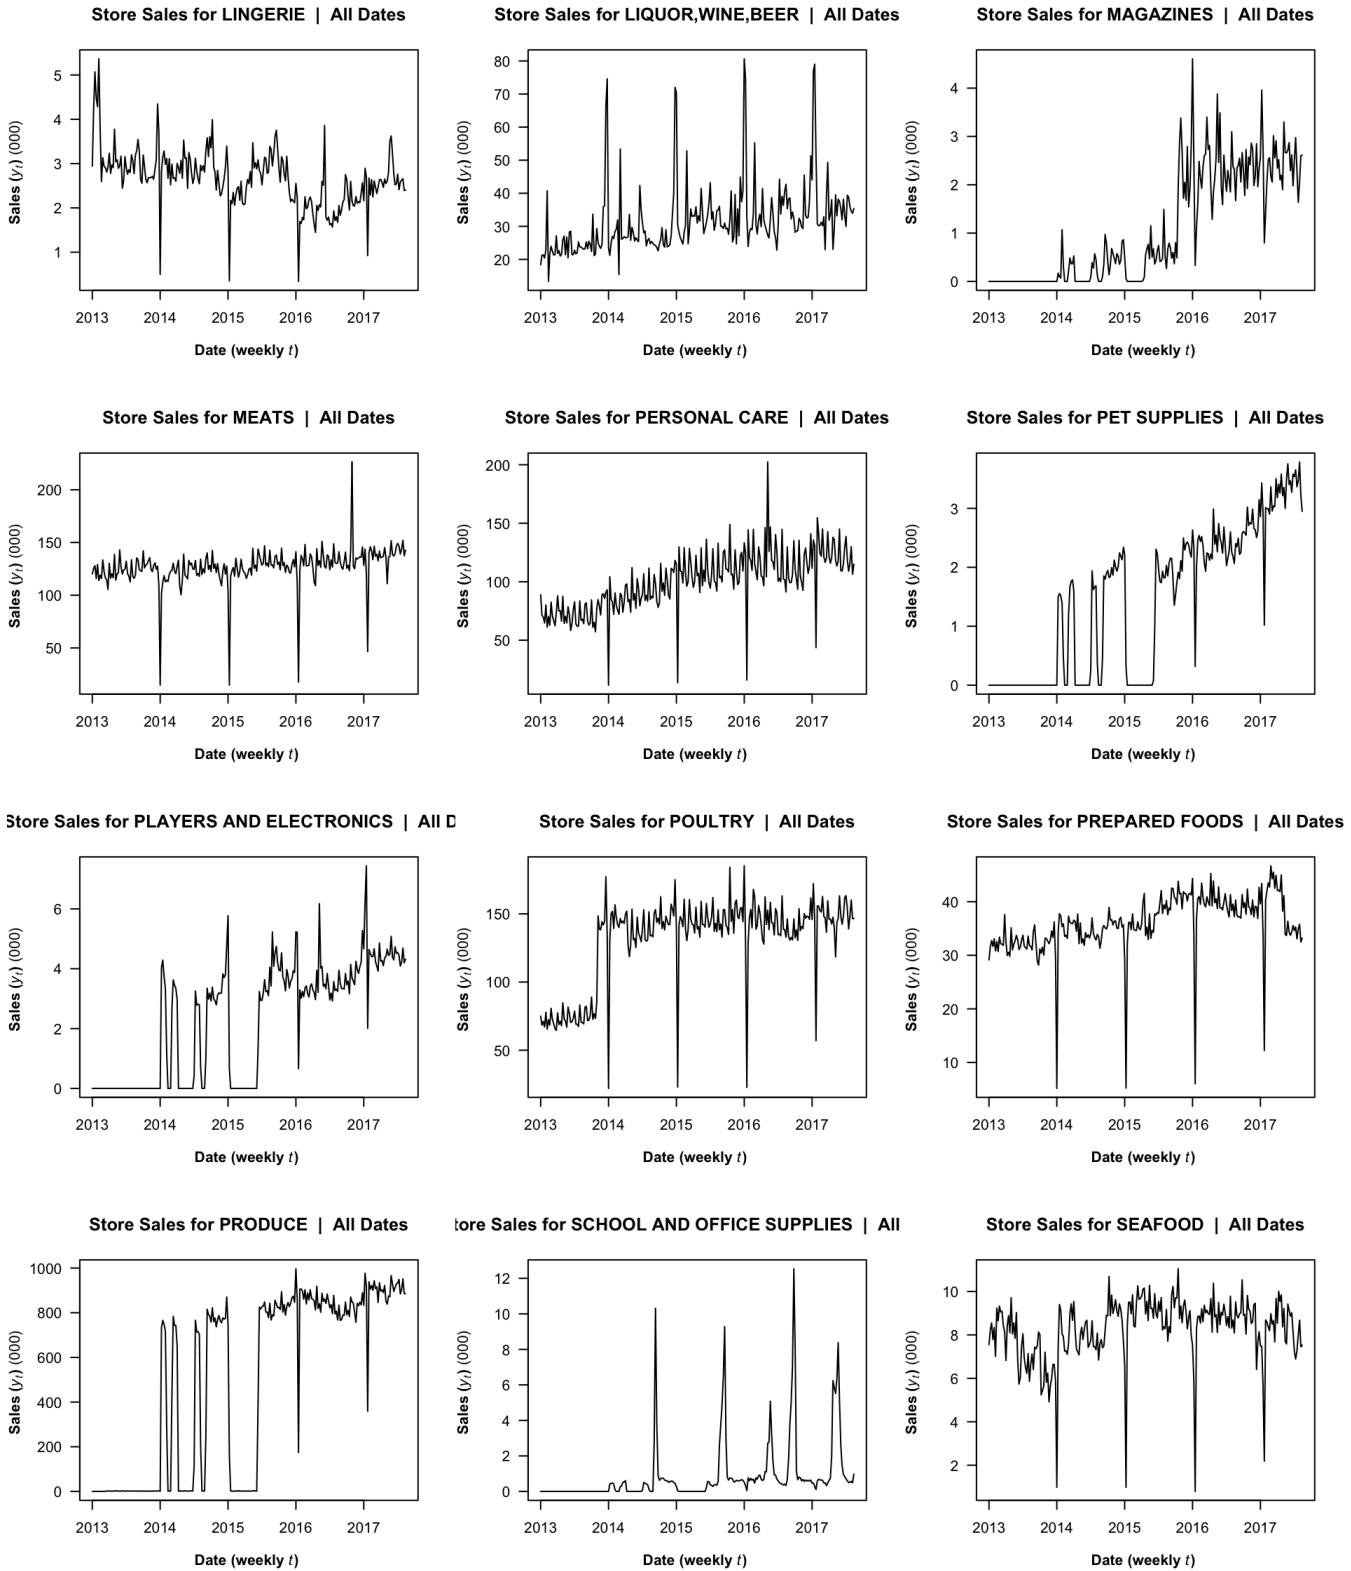
for (f in unique(sales_wk_df$family)) {
  # Subset data by product family and create time series
  df <- filter(sales_wk_df, family == f)
  df_ts <- ts(df$sales,
    start = c(sales_begin_year, sales_begin_week),
    end = c(sales_end_year, sales_end_week), freq = 52)

  # Plot time series
  plot(df_ts, type = "l",
    main = paste("Store Sales for ", f, " | All Dates", sep = ""),
    height = 0.8,
    xlab = TeX(r"\textbf{Date (weekly  $t$ )}"),
    ylab = TeX(r"\textbf{Sales ( $y_t$ ) (000)}"),
    las = 1, cex.axis = 0.7)
}

```







```
par(opar)
```

## Pre-Process Data

```

# Placeholder code to remove rows with NAs
#df %>%
#  na.omit()

# Placeholder code to remove rows with NAs in specific column
#df %>%
#  filter(!is.na(column_name))

# Placeholder code to reremove duplicates
#df %>%
#  distinct()

# Placeholder code to remove rows by index position
#df %>%
#  filter(!row_number() %in% c(1, 2, 4))

# Remove rows not considered "core"
keep_families <- c("BEAUTY", "BEVERAGES", "BREAD/BAKERY", "CLEANING", "DAIRY",
                   "DELI", "EGGS", "FROZEN FOODS", "GROCERY", "GROCERY II", "HARDWARE",
                   "LIQUOR,WINE,BEER", "MEATS", "PERSONAL CARE", "PET SUPPLIES", "POULTRY",
                   "PREPARED FOODS", "PRODUCE", "SEAFOOD")
remove_families <- c("AUTOMOTIVE", "BABY CARE", "BOOKS", "CELEBRATION", "HOME APPLIANCES",
                     "HOME AND KITCHEN I", "HOME AND KITCHEN II", "HOME CARE", "LADIESWEAR",
                     "LAWN AND GARDEN", "LINGERIE", "MAGAZINES", "PET SUPPLIES",
                     "PLAYERS AND ELECTRONICS", "SCHOOL AND OFFICE SUPPLIES")
sales_core_df <- (sales_df %>%
                    filter(family %in% keep_families))

```

```

# Mean of missing oil_df data
mean(is.na(oil_df))

```

```
[1] 0.02
```

```

# Option: Drop missing records since they only account for 2% of data
# na.omit(oil_df)
# oil_df2: New 'oil_df' data frame without missing values
# Foward and backward fill for missing values
oil_df2 <- na.locf(oil_df, fromLast = TRUE)
# Confirm oil_df2 has no more missing values
sum(is.na(oil_df2))

```

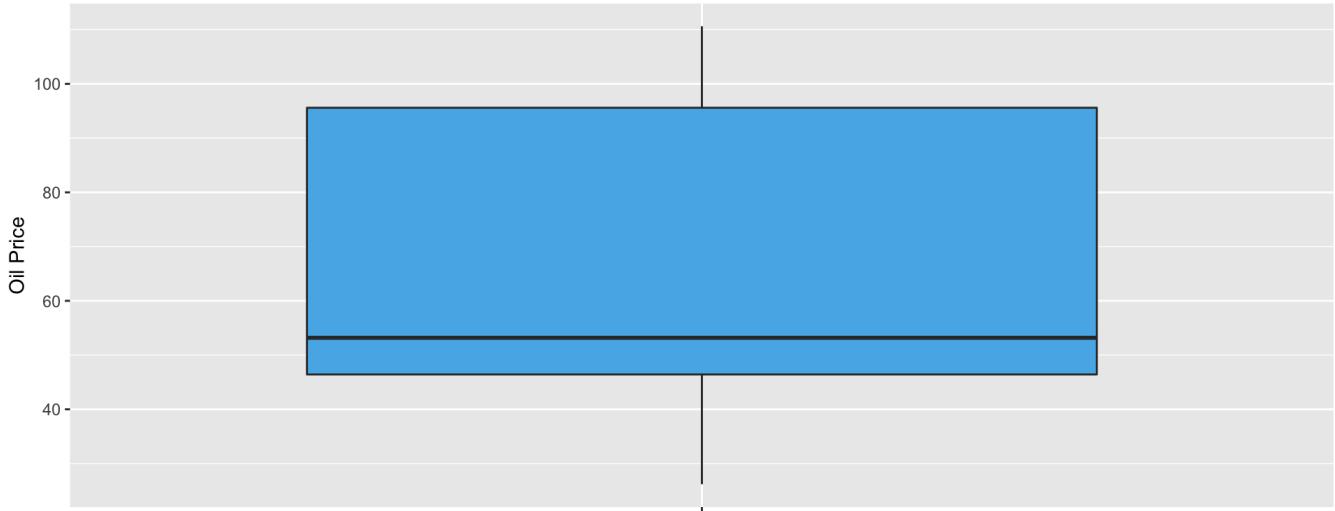
```
[1] 0
```

```

# Box plot: oil_df2
oil_outliers <- ggplot(oil_df2) +
  aes(x = "", y = dcoilwtico) +
  geom_boxplot(fill = "#56B4E9") +
  labs(title = "Oil Box Plot",
       x = "",
       y = "Oil Price",
       bty = "l")
oil_outliers

```

## Oil Box Plot



```
# Convert string mm/dd/yyyy to Date values and confirm sort
oil_df <- (oil_df %>%
  mutate(date = as.Date(date, format = "%Y-%m-%d"),) %>%
  arrange(date))

# Add oil prices to sales dataset for potential use as external predictor
sales_core_df <- left_join(sales_core_df, oil_df, by = c("date"))
```

## Series Visualization - Select ("Core") Product Families

```
# Re-aggregate base dataframe from daily to weekly for all product families
sales_wk_df <- as.data.frame(sales_core_df %>%
  mutate(year = year(date), week = week(date)) %>%
  group_by(year, week) %>%
  summarize(sales = sum(sales / 1000.0), dcoilwtico = mean(dcoilwtico)))

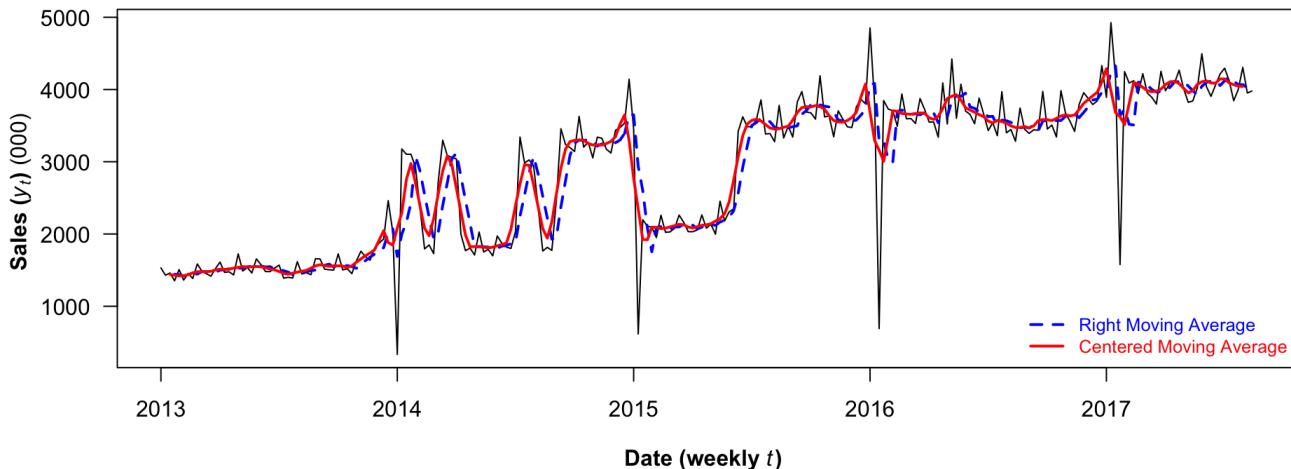
# Re-create overall time series
sales_begin_year <- head(sales_wk_df$year, 1)
sales_begin_week <- head(sales_wk_df$week, 1)
sales_end_year <- tail(sales_wk_df$year, 1)
sales_end_week <- tail(sales_wk_df$week, 1)
sales_ts <- ts(sales_wk_df$sales,
  start = c(sales_begin_year, sales_begin_week),
  end = c(sales_end_year, sales_end_week), freq = 52)

# Plot overall time series with moving averages
plot(sales_ts, type = "l",
  main = "Store Sales for Select Product Families | Weekly | All Dates",
  xlab = TeX(r"(\textbf{Date (weekly \textit{$t$})})"),
  ylab = TeX(r"(\textbf{Sales (\textit{$y_t$})} (000))"),
  las = 1, cex.axis = 0.7)

sales_l_ma <- rollmean(sales_ts, k = 4, align = "right")
sales_c_ma <- ma(sales_ts, order = 4)
lines(sales_l_ma, lwd = 2, lty = 2, col = "blue")
lines(sales_c_ma, lwd = 2, lty = 1, col = "red")

legend("bottomright",
  legend = c("Right Moving Average", "Centered Moving Average"),
  col = c("blue", "red"),
  lwd = 2, lty = c(2, 1), cex = 0.8,
  box.lty = 0, bg = NULL, text.col = c("blue", "red"))
```

## Store Sales for Select Product Families | Weekly | All Dates



## Partition Series

```
# Use one year (52 weeks) as validation period (representative set of quarters, seasons)
sales_n_valid <- 52
sales_n_train <- length(sales_ts) - sales_n_valid

# Split data into training and validation periods
sales_train_ts <- window(sales_ts, start = c(sales_begin_year, 1), end = c(sales_begin_year, sales_n_train))
sales_valid_ts <- window(sales_ts, start = c(sales_begin_year, sales_n_train + 1), end = c(sales_begin_year, sales_n_train + sales_n_valid))

# Set x axis values based on ts ranges
x_train <- sales_begin_year
x_valid <- sales_begin_year + ((sales_n_train + 1) / 52)
x_future <- x_valid + ((sales_n_valid + 1) / 52)
```

## Apply Forecasting Method(s)

```
# Initialize dataframe for tracking results across model approaches
results_df <- data.frame()
```

## Naive Forecast Baseline

```
# Fit seasonal naive model
sales_snaive_pred <- snaive(sales_train_ts, h = sales_n_valid)

# Summarize seasonal naive forecast results
accuracy(sales_snaive_pred, sales_valid_ts)
```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
Training set	784	1217	996	17	38	1.0	0.5	NA
Test set	261	752	507	4	14	0.5	-0.2	0.7

```

results_df <- rbind(results_df, data.frame(Model = "Naive (baseline)",
                                         Set = c("Train", "Test"),
                                         accuracy(sales_snaive_pred, sales_valid_ts)))

# Plot seasonal naive forecaster
ymax <- max(sales_ts) * 1.5
xmax <- x_future + 1

plot(sales_snaive_pred, type = "l",
      main = "Store Sales for Select Product Families | Seasonal Naive Forecaster",
      xlab = TeX(r"(\textbf{Date (weekly \textit{\$t\$})} )"), ylab = TeX(r"(\textbf{Sales (\textit{\$y_t\$}) (000)} )"),
      las = 1, cex.axis = 0.7,
      xlim = c(x_train, xmax), ylim = c(0, ymax))

# Add additional lines
lines(sales_valid_ts, lwd = 2, col = "red")

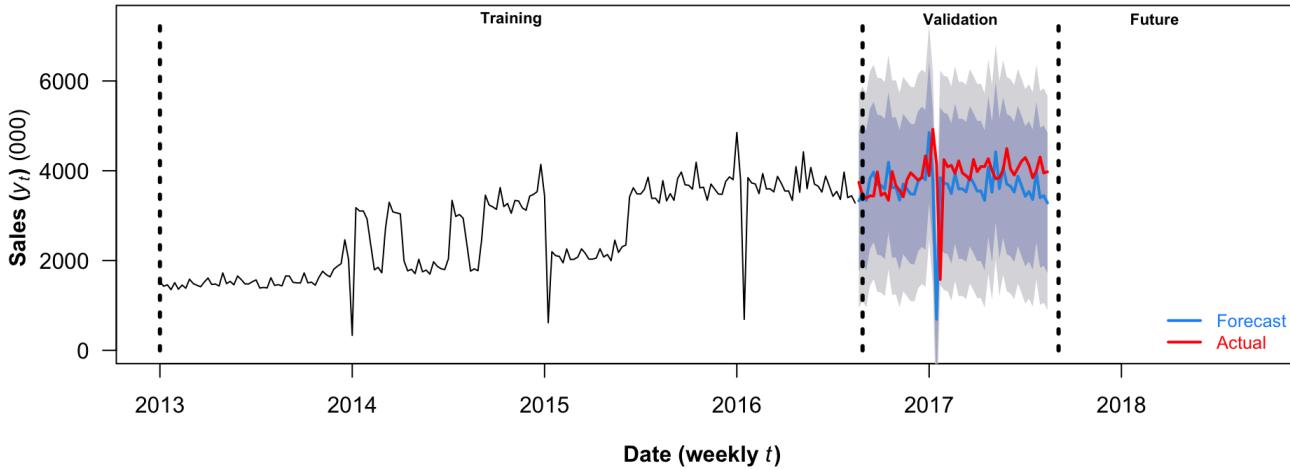
# Add plot definitions
lines(c(x_train, x_train), c(0, ymax), lwd = 3, lty = 3)
lines(c(x_valid, x_valid), c(0, ymax), lwd = 3, lty = 3)
lines(c(x_future, x_future), c(0, ymax), lwd = 3, lty = 3)

text(x_train + ((x_valid - x_train) * 0.5), ymax, "Training", font = 2, cex = 0.7)
text(x_valid + ((x_future - x_valid) * 0.5), ymax, "Validation", font = 2, cex = 0.7)
text(x_future + 0.5, ymax, "Future", font = 2, cex = 0.7)

legend("bottomright",
       legend = c("Forecast", "Actual"),
       col = c("dodgerblue", "red"),
       lwd = 2, lty = 1, cex = 0.8,
       box.lty = 0, bg = NULL, text.col = c("dodgerblue", "red"))

```

### Store Sales for Select Product Families | Seasonal Naive Forecaster



### Data-Driven: Smoothing Methods

```

# Fit Holt-Winter's model with (M)ultiplicative error, (A)dditive trend, and (M)ultiplicative seasonality
#   for last year of training series
sales_hw <- ets(window(sales_train_ts, start = c(sales_begin_year, sales_n_train + 1 - sales_n_valid),
                      end = c(sales_begin_year, sales_n_train + 1)),
                  model = "MAM")

# Forecast (predict) for validation period
sales_hw_pred <- forecast(sales_hw, h = length(sales_valid_ts), level = 0)

# Summarize forecast (validation) performance
accuracy(sales_hw_pred, sales_valid_ts)

```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
Training set	-39	523	258	-9	14	NaN	-0.002	NA
Test set	578	781	673	13	18	NaN	0.334	0.8

```

results_df <- rbind(results_df, data.frame(Model = "Holt-Winter's Smoothing",
                                         Set = c("Train", "Test"),
                                         accuracy(sales_hw_pred, sales_valid_ts)))

# Plot Holt-Winter's model
plot(sales_hw_pred, type = "l",
      main = "Store Sales for Select Product Families | Holt-Winter's Forecaster",
      xlab = TeX(r"(\textbf{Date (weekly } \textit{t} \text{)})"),
      ylab = TeX(r"(\textbf{Sales (\textit{y}_t)}) (000)"),
      las = 1, cex.axis = 0.7,
      xlim = c(x_train, xmax), ylim = c(0, ymax))

# Add additional lines
lines(sales_hw$fitted, lwd = 3, col = adjustcolor("dodgerblue", alpha.f = 0.3))
lines(sales_valid_ts, lwd = 2, col = "red")

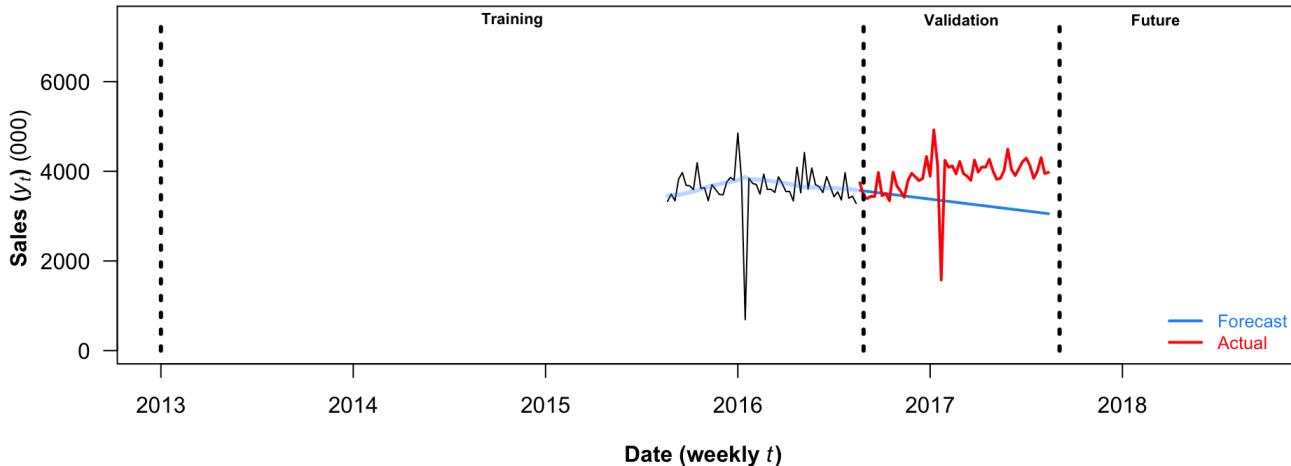
# Add plot definitions
lines(c(x_train, x_train), c(0, ymax), lwd = 3, lty = 3)
lines(c(x_valid, x_valid), c(0, ymax), lwd = 3, lty = 3)
lines(c(x_future, x_future), c(0, ymax), lwd = 3, lty = 3)

text(x_train + ((x_valid - x_train) * 0.5), ymax, "Training", font = 2, cex = 0.7)
text(x_valid + ((x_future - x_valid) * 0.5), ymax, "Validation", font = 2, cex = 0.7)
text(x_future + 0.5, ymax, "Future", font = 2, cex = 0.7)

legend("bottomright",
       legend = c("Forecast", "Actual"),
       col = c("dodgerblue", "red"),
       lwd = 2, lty = 1, cex = 0.8,
       box.lty = 0, bg = NULL, text.col = c("dodgerblue", "red"))

```

### Store Sales for Select Product Families | Holt-Winter's Forecaster



### Model-Driven: Regression

```

# Fit basic regression model with trend and seasonality
sales_lm <- tslm(sales_train_ts ~ trend + I(trend ^ 2) + season)

# Forecast (predict) for validation period
sales_lm_pred <- forecast(sales_lm, h = sales_n_valid, level = 0)

# Summarize performance metrics
accuracy(sales_lm_pred, sales_valid_ts)

```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
Training set	-0.0000000000000005	525	393	-8	21	0.4	0.47	NA
Test set	-529.415459820942715	735	583	-16	17	0.6	0.06	0.7

```

results_df <- rbind(results_df, data.frame(Model = "Simple Regression",
                                         Set = c("Train", "Test"),
                                         accuracy(sales_lm_pred, sales_valid_ts)))

# Plot simple regression model
plot(sales_lm_pred, type = "l",
      main = "Store Sales for Select Product Families | Regression Forecaster",
      xlab = TeX(r"(\textbf{Date (weekly } t\textbf{)})"),
      ylab = TeX(r"(\textbf{Sales (\textit{$y_t$}) (000)})"),
      las = 1, cex.axis = 0.7,
      xlim = c(x_train, xmax), ylim = c(0, ymax))

# Add additional lines
lines(sales_lm$fitted, lwd = 3, col = adjustcolor("dodgerblue", alpha.f = 0.3))
lines(sales_valid_ts, lwd = 2, col = "red")

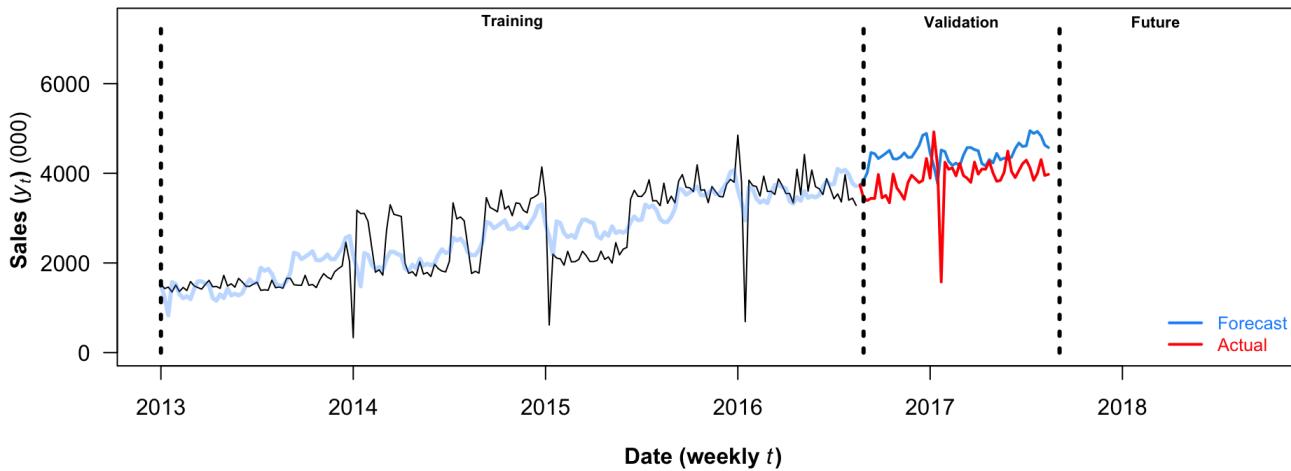
# Add plot definitions
lines(c(x_train, x_train), c(0, ymax), lwd = 3, lty = 3)
lines(c(x_valid, x_valid), c(0, ymax), lwd = 3, lty = 3)
lines(c(x_future, x_future), c(0, ymax), lwd = 3, lty = 3)

text(x_train + ((x_valid - x_train) * 0.5), ymax, "Training", font = 2, cex = 0.7)
text(x_valid + ((x_future - x_valid) * 0.5), ymax, "Validation", font = 2, cex = 0.7)
text(x_future + 0.5, ymax, "Future", font = 2, cex = 0.7)

legend("bottomright",
       legend = c("Forecast", "Actual"),
       col = c("dodgerblue", "red"),
       lwd = 2, lty = 1, cex = 0.8,
       box.lty = 0, bg = NULL, text.col = c("dodgerblue", "red"))

```

### Store Sales for Select Product Families | Regression Forecaster



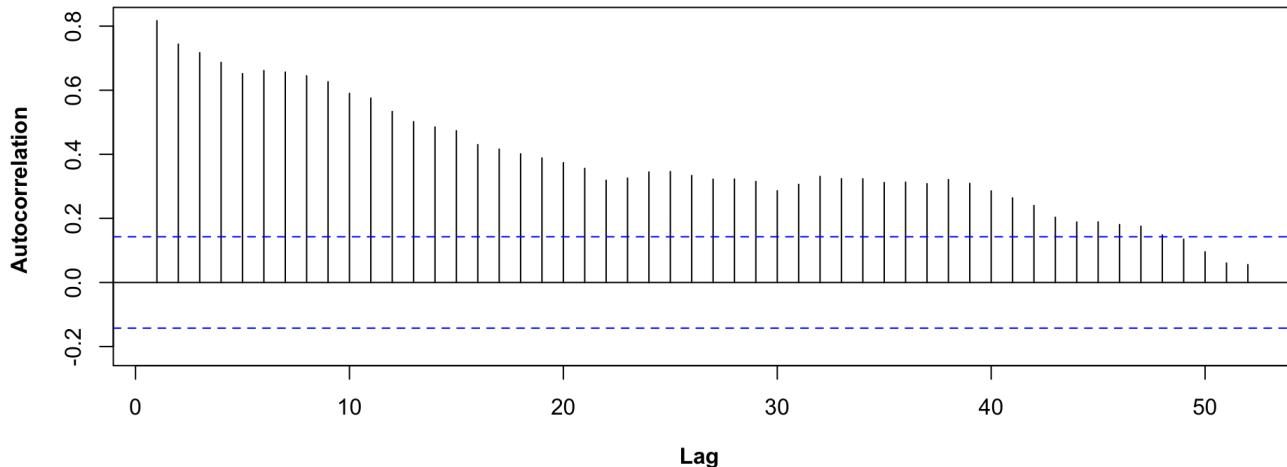
### Model-Driven: Regression w/Autocorrelation

```

# Compute and plot autocorrelation
Acf(sales_train_ts, lag.max = 52, type = "correlation", plot = TRUE,
     main = "Store Sales for Select Product Families | Autocorrelation at Different Lags",
     xlab = TeX(r"(\textbf{Lag})"),
     ylab = TeX(r"(\textbf{Autocorrelation})"),
     xaxt = "n",
     level = 95)

```

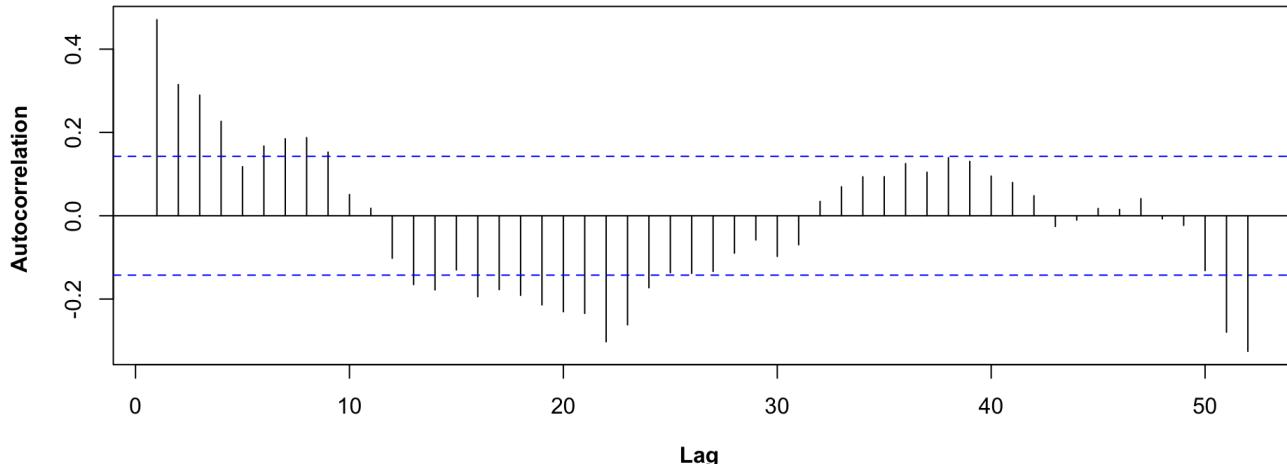
## Store Sales for Select Product Families | Autocorrelation at Different Lags



## Model-Driven: Regression w/Autocorrelation - AR(1)

```
# Compute and plot autocorrelation for regression residuals
Acf(sales_lm$residuals, lag.max = 52, type = "correlation", plot = TRUE,
  main = "Regression Residuals | Autocorrelation at Different Lags",
  xlab = TeX(r"\textbf{Lag}"), ylab = TeX(r"\textbf{Autocorrelation}")),
  xaxt = "n",
  level = 95)
```

## Regression Residuals | Autocorrelation at Different Lags



```
# Fit AR(1) second-layer model to regression residuals and forecast (predict) vs. validation;
#   use AR of 1, MA of 0, and integration (differencing) of 0
sales_res_arima <- Arima(sales_lm$residuals, order = c(1, 0, 2))
sales_res_pred <- forecast(sales_res_arima, h = sales_n_valid)

# Summarize forecast (validation) performance
accuracy(sales_res_pred, sales_valid_ts)
```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
Training set	-1	455	288	278	328	0.4	-0.001	NA
Test set	3926	3949	3926	101	101	5.1	0.015	4

```

results_df <- rbind(results_df, data.frame(Model = "AR(1)",
                                         Set = c("Train", "Test"),
                                         accuracy(sales_res_pred, sales_valid_ts)))

rymin <- min(sales_res_pred$residuals)
rymax <- max(sales_res_pred$residuals) * 1.5

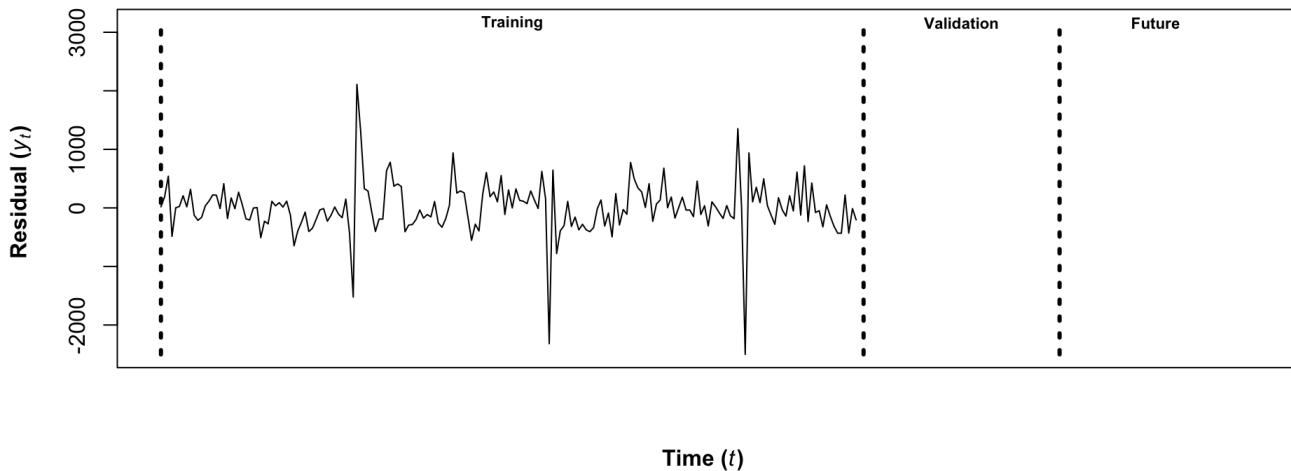
# Plot residuals
plot(sales_res_pred$residuals, type = "l",
      main = "Store Sales for Select Product Families | Residuals (errors) from AR(1)",
      xlab = TeX(r"(\textbf{Time} (\textit{$t$}))"),
      ylab = TeX(r"(\textbf{Residual} (\textit{$y_t$}))"),
      xaxt = "n",
      xlim = c(x_train, x_max), ylim = c(rymin, rymax))

# Add plot definitions
lines(c(x_train, x_train), c(rymin, rymax), lwd = 3, lty = 3)
lines(c(x_valid, x_valid), c(rymin, rymax), lwd = 3, lty = 3)
lines(c(x_future, x_future), c(rymin, rymax), lwd = 3, lty = 3)

text(x_train + ((x_valid - x_train) * 0.5), rymax, "Training", font = 2, cex = 0.7)
text(x_valid + ((x_future - x_valid) * 0.5), rymax, "Validation", font = 2, cex = 0.7)
text(x_future + 0.5, rymax, "Future", font = 2, cex = 0.7)

```

**Store Sales for Select Product Families | Residuals (errors) from AR(1)**

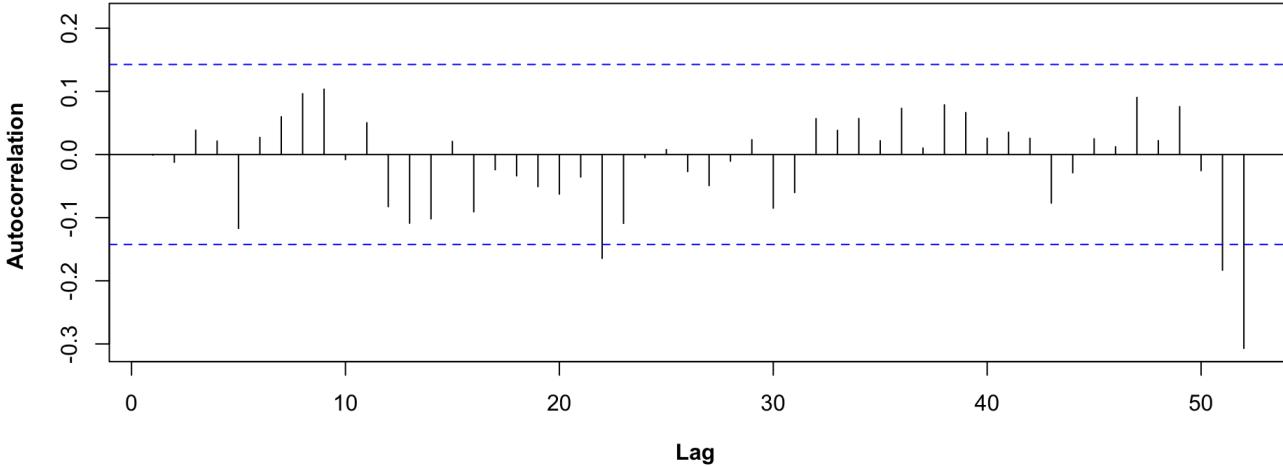


```

# Compute and plot autocorrelation for residual of residuals
Acf(sales_res_arima$residuals, lag.max = 52, type = "correlation", plot = TRUE,
      main = "AR(1) Residuals | Autocorrelation at Different Lags",
      xlab = TeX(r"(\textbf{Lag})"),
      ylab = TeX(r"(\textbf{Autocorrelation})"),
      xaxt = "n",
      level = 95)

```

## AR(1) Residuals | Autocorrelation at Different Lags



## Model-Driven: Regression w/Autocorrelation - ARIMA

```
# Fit ARIMA model
sales_arima <- Arima(sales_train_ts, order = c(1, 0, 0), seasonal = c(0, 0, 2))
sales_arima_pred <- forecast(sales_arima, h = length(sales_valid_ts))

# Summarize forecast (validation) performance
accuracy(sales_arima_pred, sales_valid_ts)
```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
Training set	-20	490	303	-8	17	0.3	-0.1	NA
Test set	1582	1680	1607	39	41	1.6	0.5	2

```
results_df <- rbind(results_df, data.frame(Model = "ARIMA",
                                             Set = c("Train", "Test"),
                                             accuracy(sales_arima_pred, sales_valid_ts)))

# Plot model
plot(sales_arima_pred, type = "l",
      main = "Store Sales for Select Product Families | ARIMA Forecaster",
      xlab = TeX(r"(\textbf{Date (weekly \textit{\$t\$})})"),
      ylab = TeX(r"(\textbf{Sales (\textit{\$y_t\$})} (000))"),
      las = 1, cex.axis = 0.7,
      xlim = c(x_train, xmax), ylim = c(0, ymax))

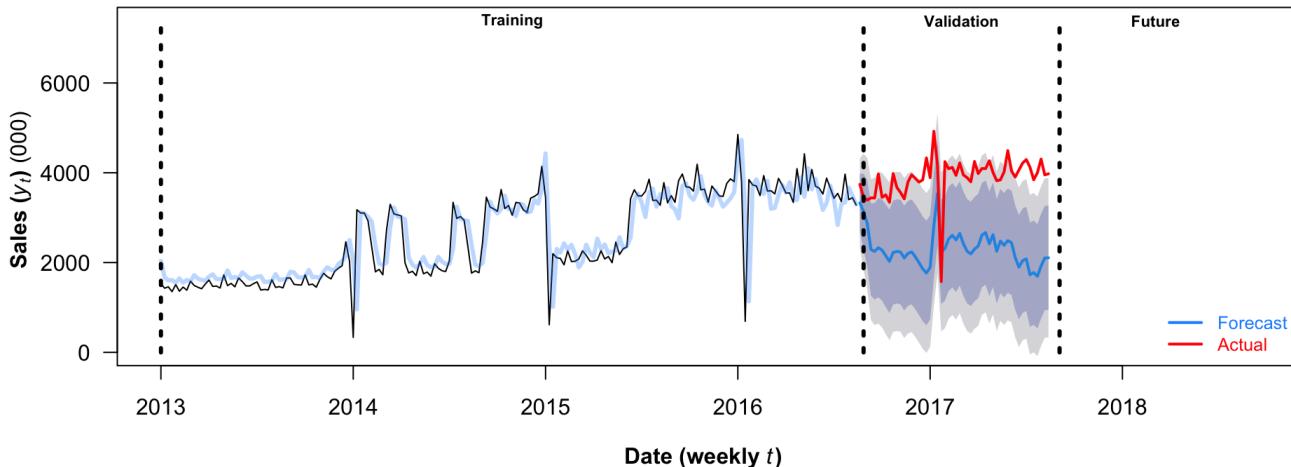
# Add additional lines
lines(sales_arima$fitted, lwd = 3, col = adjustcolor("dodgerblue", alpha.f = 0.3))
lines(sales_valid_ts, lwd = 2, col = "red")

# Add plot definitions
lines(c(x_train, x_train), c(0, ymax), lwd = 3, lty = 3)
lines(c(x_valid, x_valid), c(0, ymax), lwd = 3, lty = 3)
lines(c(x_future, x_future), c(0, ymax), lwd = 3, lty = 3)

text(x_train + ((x_valid - x_train) * 0.5), ymax, "Training", font = 2, cex = 0.7)
text(x_valid + ((x_future - x_valid) * 0.5), ymax, "Validation", font = 2, cex = 0.7)
text(x_future + 0.5, ymax, "Future", font = 2, cex = 0.7)

legend("bottomright",
       legend = c("Forecast", "Actual"),
       col = c("dodgerblue", "red"),
       lwd = 2, lty = 1, cex = 0.8,
       box.lty = 0, bg = NULL, text.col = c("dodgerblue", "red"))
```

## Store Sales for Select Product Families | ARIMA Forecaster



### Model-Driven: Regression w/External Information

```
#*****CODE HERE*****
```

### Model-Driven: Regression w/Autocorrelation - ARIMA + External Predictor(r)

```
# Identify outcome (target) - sales
sales_train_y <- ts(sales_wk_df$sales, freq = 52)

# Identify external predictor(s) - oil price
sales_X <- c("dcoilwtico")
sales_train_X <- as.matrix(sales_wk_df[, sales_X])

# Fit using auto.arima for sense of best component configuration, and summarize results both without and with external
# predictor(s)
sales_aam <- auto.arima(sales_train_y)
sales_aamX <- auto.arima(sales_train_y, xreg = sales_train_X)
summary(sales_aam)
```

```
Series: sales_train_y
ARIMA(0,1,2)(0,0,1)[52]

Coefficients:
      m1      m2     smal
      -0.42   -0.30   -0.31
  s.e.  0.07   0.07   0.09

sigma^2 = 255333: log likelihood = -1867
AIC=3741   AIcC=3741   BIC=3755

Training set error measures:
      ME RMSE MAE MPE MAPE MASE    ACF1
Training set 38 501 309 -4   15  0.4 -0.003
```

```
summary(sales_aamX)
```

```

Series: sales_train_y
Regression with ARIMA(0,0,0) errors

Coefficients:
intercept xreg
921 -5.9
s.e. 14 0.2

sigma^2 = 1.1: log likelihood = -11
AIC=28 AICc=28 BIC=36

Training set error measures:
ME RMSE MAE MPE MAPE MASE ACF1
Training set -0.000000000000001 9 9 -0.1 1 0.01 NA

```

```

# Fit ARIMA model
*****NOTE HAVE BUG IN ARIMA W/EXTERNAL PREDICTOR*****
sales_am <- Arima(sales_train_y, order = c(0, 0, 1))
#sales_amX <- Arima(sales_train_y, order = c(0, 0, 1), xreg = sales_train_X)
summary(sales_am)

```

```

Series: sales_train_y
ARIMA(0,0,1) with non-zero mean

Coefficients:
ma1 mean
0.63 2878
s.e. 0.04 79

sigma^2 = 581675: log likelihood = -1973
AIC=3952 AICc=3952 BIC=3962

Training set error measures:
ME RMSE MAE MPE MAPE MASE ACF1
Training set 1 760 635 -14 31 0.7 0.3

```

```
#summary(sales_amX)
```

## Evaluate & Compare Performance

```

# Show table of summary results
print(results_df, row.names = FALSE)

```

Model	Set	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil.s.U
Naive (baseline)	Train	783.614908859564025	1217	996	17	38	1.0	0.477	NA
Naive (baseline)	Test	261.373668023741175	752	507	4	14	0.5	-0.237	0.7
Holt-Winter's Smoothing	Train	-38.619403647477426	523	258	-9	14	NaN	-0.002	NA
Holt-Winter's Smoothing	Test	578.171434927598284	781	673	13	18	NaN	0.334	0.8
Simple Regression	Train	-0.000000000000005	525	393	-8	21	0.4	0.471	NA
Simple Regression	Test	-529.415459820942715	735	583	-16	17	0.6	0.056	0.7
AR(1)	Train	-1.225461539994329	455	288	278	328	0.4	-0.001	NA
AR(1)	Test	3925.739961464764747	3949	3926	101	101	5.1	0.015	4.0
ARIMA	Train	-20.344171628362030	490	303	-8	17	0.3	-0.149	NA
ARIMA	Test	1581.833441101568951	1680	1607	39	41	1.6	0.516	1.7

```

# Visualize model comparison
*****CODE HERE*****

```

## Implement Forecasts/System

<sup>1</sup>.University of San Diego, cvu@sandiego.edu (mailto:cvu@sandiego.edu) ↵

<sup>2</sup>.University of San Diego, dfriesen@sandiego.edu (mailto:dfriesen@sandiego.edu) ↵